

What is FIS-B 978?

FIS-B 978 ('fisb-978') is a set of programs that demodulates and error corrects FIS-B (Flight Information System - Broadcast) and ADS-B (Automatic Dependent Surveillance - Broadcast) messages transmitted on 978 Mhz, mostly in the United States.

It performs the same function as [FlightAware's 978](#), but is spread out over several programs that are piped together. It's main advantage over FlightAware's 978 and other 'dump-978' clones is that for moderate to weaker signals it can provide much higher packet decode rates, since it takes into account that samples at the Nyquist sampling rate may not be optimal and will shift them to more optimal levels.

For strong signals, 'dump-978' and 'fisb-978' will perform similarly.

'fisb-978' is composed of two main parts and one optional part:

- **demod_978** is a C language program which takes raw data from an SDR program at 978 Mhz and sample rate of 2.083334 Mhz (with 16-bit complex integer output (CS16)) and will demodulate the data, detect sync words, and output packets along with attributes of the packet to standard output. Typically, 'ec-978.py' will read this data.
- **ec_978.py** takes the data from 'demod_978' and will use Reed-Solomon error correction to produce an output packet as hex bytes (similar to 'dump-978'). It uses a number of techniques to provide a higher packet decode rate. The output is sent to standard output, where it can be used as is, or processed by 'server_978'.
- **server_978.py** takes the output from 'ec_978.py' and will provide a TCP server where clients can connect and receive the data. It is optional if you don't want to serve the data via TCP.

The program is broken into parts for modularity and speed. C is very fast for searching through the input stream and finding sync packets (numpy is great at the demodulation part, but horrid at searching for sync codes). Python using numpy (and a Reed-Solomon library linked to a C library), is very quick at error correcting the packets and manipulating them to provide more optimum bit levels for better error correction rates.

A few things to consider before using:

- Since 'demod_978.c' uses type-punning, a *compiler that is friendly to that is required*. GCC is such a compiler. *All code expects little-endian byte order*. This will work on most common architectures in use today. If needed, big-endian can be added as a future feature.
- 'server_978.py' uses a `select()` statement using both sockets and file I/O. As such, *this will usually not work on Windows* (it should work fine if you are using the 'Linux Subsystem for Linux').
- 'fisb-978' may improve the number of packets you can error correct, but it should be low on your list of improvements that actually matter. Nothing beats a good antenna, and a good 978MHz LNA and filter. Also, better quality radios are always a plus.

FIS-B 978 is not designed, nor intended, to be used for any purpose other than fun exploration.

Warning: FIS-B 978 is **NOT** intended for actual flight use, nor to be used as a component of a system for any flight related operations (flight planning, etc). It is strictly a fun hobby program with many bugs. It is based loosely on draft standards, not actual standards, and has not undergone any formal testing. **DO NOT USE OTHER THAN FOR FUN- NOT FLIGHT!!**

Getting things running

There is not much setup to do. For `decode_978` just cd to the cloned directory and type `make`. You should see something like:

```
$ make  
gcc -c -o demod_978.o demod_978.c -I. -O3 -Wall -funroll-loops  
gcc -o demod_978 demod_978.o -I. -O3 -Wall -funroll-loops
```

There is nothing to do for `server_978.py`. It should work out of the box.

`ec_978.py` requires *numpy*, a C library, and the Reed-Solomon interface to the C library.

There are lots of ways to install *numpy*. Go to numpy.org if you need help. If you have `pip3`, just type:

```
pip3 install numpy
```

If you going to use `eye.py` to generate eye diagrams, you will need to install *scipy*, *matplotlib*, and *colour* (as well as *numpy*). This can all be done with:

```
pip3 install numpy matplotlib scipy colour
```

For Reed-Solomon, visit <https://pypi.org/project/pyreedsolomon/>. The best way to install this is **not** with `pip3`. Make a clone as follows:

```
# Do the following line only if you have it already installed  
# or think you might.  
sudo pip3 uninstall pyreedsolomon  
  
git clone --recursive https://github.com/epeters13/pyreedsolomon.git  
cd pyreedsolomon  
sudo python3 setup.py install  
sudo ldconfig
```

`ec_978.py` should now have all the prerequisites installed.

Initial Usage

To use, you just pipe the output from your SDR program to ‘`demod_978`’ and then to ‘`ec_978.py`’. Optionally, you can pipe it to ‘`server_978.py`’ if you want to distribute it. Not including the server will send the output to a terminal, which is a good place to start.

Here is an example of a script I use for normal decoding (based on the *SDRplay RSP1A*). It can be found in `scripts/sdrplay-demod`:

```
# Script used for using rx_sdr with sdrplay to capture raw data  
# and send them to demod_978, ec_978, and server_978.  
#  
# You can change the first portion to reflect your SDR radio, but use  
# the following settings:  
# frequency: 978 MHz
```

```

# sample rate:          2083334
# output type:         Complex 16 bit integer (usually called CS16)
#
# If you just want to see decoded output printed, leave off
# '| ./server_978.py'.
#
rx_sdr -d driver=sdrplay -t biasT_ctrl=true,rfgain_sel=1 -g 25 -F CS16 \
-f 978000000 -s 2083334 - | ./demod_978 | ./ec_978.py | ./server_978.py

```

Leave off 'server_978.py' when starting, so you can see the output on the terminal.

You will need to substitute your SDR program and settings. The settings must include the frequency, sample rate, and output type (CS16) shown above. You just pipe the raw output through 'demod_978' and 'ec_978.py'. This will give you the decoded hex strings for FIS-B and ADS-B. To serve it remotely, pipe that output to 'server_978.py'.

For an 8-bit *RTLSDR* or *RadarBox 978 FlightStick*, the above command can be modified as (also found in `scripts/rtlsdr-demod`):

```

rx_sdr -d driver=rtlsdr -F CS16 -g 40 -f 978000000 -s 2083334 - \
| ./demod_978 | ./ec_978.py | ./server_978.py

```

You won't get the performance out of an 8-bit SDR as you will with something with a higher bit ADC. Also, for all radios, a good filtered preamp is the next most important thing after a good antenna. I have found the [Uputronics 978MHz UAT filtered preamp](#) to be an excellent performer. In the United States you can get them at [AIRSPY.us](#) (disclosure: I am not sponsored by any product, nor do I have affiliate links).

When first starting, it is a good idea to explore your setup to find the best level (minimum signal strength) for 'demod_978'. By default, 'demod_978' is tuned for FIS-B and not ADS-B. It will not attempt to filter packets below a certain level. This is because if the level is set too low, you will get a huge number of noise packets that match the sync codes, but are just noise. This minimum level will be lower for ADS-B than for FIS-B.

To see what levels your packets are successfully decoding at, set the minimum level in 'demod_978' to 0 with the `-l 0` argument, and turn on the *lowest levels* flag in 'ec_978.py' with `--ll`. The `--ll` argument will show the lowest level received for FIS-B, ADS-B short, and ADS-B long. The output is sent to standard error, so best to pipe standard output to `/dev/null` so it doesn't get lost in decoded packets. For example:

```

<your SDR program> | ./demod_978 -l 0 | ./ec_978.py --ll >/dev/null
lowest ADS-B (L) signal: 1.01
lowest ADS-B (S) signal: 0.63
lowest FIS-B      signal: 15.66
lowest FIS-B      signal: 14.52
lowest FIS-B      signal: 5.69
lowest FIS-B      signal: 5.66
lowest FIS-B      signal: 3.77
lowest FIS-B      signal: 3.48
lowest FIS-B      signal: 2.23
lowest ADS-B (S) signal: 0.28
lowest ADS-B (S) signal: 0.17
lowest FIS-B      signal: 1.93

```

In this example, the lowest FIS-B packet was at signal strength 1.93 and the lowest ADS-B packet was at 0.17. So if you are only interested in FIS-B, the default level of 0.9 should be fine. For ADS-B, something like 0.1 might work. This can be set in 'demod_978' with the `-l` argument.

It should be noted that the Reed-Solomon error correction in ADS-B and FIS-B is not a guarantee that a message was decoded *correctly*. Reed-Solomon has a number of parity bytes (for our purposes, these are bytes, not bits). So if you have 14 bytes of parity (like in an ADS-B long message), that means it will detect and correct up to 7 error bytes that it finds, or 14 that you know exist (erasures). Or some combination of the above. *If the message actually has more errors than this, all bets are off, and Reed-Solomon may declare that the message is fine when it isn't.* None of these messages have a CRC code, or other error detection mechanism to double check that a decode is correct. So if you decode random noise, it is very possible to get Reed-Solomon certified garbage packets.

In FIS-B, there is an imperfect mechanism you can use to double check for a correct message. It isn't perfect, but it is better than nothing. In FIS-B, each message is made up of UAT Frames. Each frame comes with a length, and UAT Frames are chained together until the end of the total message. Either the frames will fit exactly in a message, or they will end before the end of the message. In which case the remainder of the message will be zeros. If a message doesn't pass this test, it isn't valid. Passing this test doesn't mean it's valid either, but the likelihood is much higher. Currently, because essentially all of the decodes from `ec_978.py` are correct, this mechanism is not used as a double check.

Explanation of program output

Output from `ec_978.py` will show three types of packets: FIS-B, ADS-B long, and ADS-B short. A FIS-B packet will look like:

Note that in reality, all packets are a single line. The examples are broken up for clarity. The '+' at the beginning indicates a FIS-B packet. ADS-B packets start with '-'.

The actual message is the set of hex characters that follow

After that, separated by ‘‘::’’, are three items:

- `rs=` is about Reed Solomon error correction and how many sync bits in the sync preamble didn't match.

In this case `rs=0/` means that all bits in the sync word matched. This number will be from 0 to 4. Each FIS-B and ADS-B packet is preceded by a 36-bit sync word. 32 of those bits have to be correct for the sync word to be considered a match.

`01:02:01:01:00:03` represents the number of Reed-Solomon errors corrected in each FIS-B block. FIS-B messages are made up of six parts (blocks), each with their own set of error correction bits. Each

block can have up to 10 errors before it is considered uncorrectable. If a packet has more than 10 errors, the number of errors will be listed as 98. You will see this when printing errors. More commonly, you will see something like: 04:99:99:99:99:99. This implies that the first block of the packet was decoded with 4 errors, but the message only had a small amount of content. This is what you will see with an empty packet (they only contain basically the latitude and longitude of the ground station). `ec_978.py` looked at the packet and determined, by only looking at block 0, that this packet is empty. In that case it doesn't even look at the other blocks. The '99' means that a packet was not even looked at.

- `ss=3.76/-20.1` represents two types of signal strength. The first number, `3.76` has no units and isn't related to anything. It is just a relative indication of the signal strength of the sync word. When data is read by `demod_978`, the demodulated data is a set of signed integers. The program keeps a running average of the absolute value of the last 72 bits (i.e. a 36 bit sync word with a set of bits in between because we are sampling at two samples per bit). In order for the program to even check for sync, this value must be larger than some threshold. By default this is 0.9, but it can be changed with the `-l` argument in `demod_978`. If the running average is above 0.9, we will attempt to match a sync word, and if we do, this value is recorded as the signal strength.

The second number, `-20.1`, is the Received Signal Strength Indication (RSSI). It is included solely for those programs that are expecting RSSI values similar to those generated by `dump978-fa` (there is a `--d978fa` switch for `ec_978.py` that will produce output in the exact format used by `dump978-fa`). It is calculated in the same way that `dump978-fa` calculates the value (which is basically the usual formula of $10 * \log_{10}(\text{power})$ plus a scalar divider that approximates the values out of `dump978-fa`). This is a negative number that gets closer to zero with a stronger signal, and further away with a weaker signal. This number is calculated using the power from the sync word.

For comparing signal strengths in FISB-978, the first number is the preferred value to use. It's linear and based on the absolute amplitude of the sync word.

- `t=1639224615.144` is the time in UTC seconds past Linux EPOCH with the number of milliseconds attached. This value is calculated as follows: `demod_978` records the time every time it reads from the disk (that usually happens 10 times a second). Whenever we decode a sync word, we calculate the time by adding 0.48 microseconds per sample for each bit from the time the disk was read, minus $0.48 * 72$ bits so the time is reflected back to when the sync word was started.

So why do we care about message arrival times? There are two primary reasons. First, having an arrival time allows you to replay messages. The most important reason for this are the FIS-B standard tests which are nothing more than packet replays. The [fisb-decode](#) repository has programs that will essentially turn time back to the starting packet time, and then wait in real time until the next packet would have been sent, and send it at the correct time relative to the original. The second reason is that the FAA, probably dating from the time reports were sent over low speed teletype circuits, uses only partial dates in reports. It is rare to get the year, month, day, and clock time, in a single message. Often, you have to use the context of when the message was sent to guess at the actual date-time. The 'fisb-decode' repository makes heavy use of message times to create JSON messages that have a complete ISO-8601 timestamp.

Another interesting thing about messages in FIS-B is that they are only sent at specific times. There are 32 channels for sending FIS-B messages. Each channel has a specific time the message will begin to be transmitted. Each second the messages for a specific ground station will be sent in 2-4 different channels (depending on ground station strength), but these channels can be predicted if you know the correct UTC second. Even stranger is that the FIS-B standard was designed such that if GPS goes out, you can use the time delay from when a ground station sent a message, to the time you received it, and calculate the

distance from you to the station (known as *pseudoranging*). Combine this with distances from other stations and you can calculate an approximate location. This concept is part of an FAA initiative known as *APNT (Alternate Positioning, Navigation, and Timing)* [And yes, if you think this through, there are lots of problems with UAT location finding, which is why you probably never heard of it].

Failed FIS-B and ADS-B messages will look something like:

```
#FAILED-FIS-B 1/99:08:99:08:10:99 ss=1.66 t=1639224737.098  
1639224737.098328.F.01663540.1  
  
#FAILED-ADS-B 2/99 ss=1.66 t=1639229048.478  
1639229048.478000.A.01658613.2
```

Again, each will be a single line. The # in the front indicates a comment. `fisb-decode` will ignore this. The data on the line is similar to what we just discussed. The last portion of the line is the attribute string that `demod_978` passed to `ec_978.py` and is used as part of the filename in case errors are being saved for further study. In order to get failed error messages, you must supply the `--ff` (FIS-B) or `--fa` (ADS-B) arguments to `ec_978.py`.

A long ADS-B message will look like:

```
-0b28c0ee3879938546c605d6100600c01105eded2ded2d0ad274030000000000000000  
;rs=0/1;ss=3.29;t=1639226996.293
```

It starts with a dash. The format is similar to FIS-B except the `rs=0/1` reflects 0 sync code errors (as in FIS-B), and one Reed-Solomon error was corrected. There is only one Reed-Solomon block in ADS-B messages, so you will only see a single number. Not six as in FIS-B.

A short ADS-B message is just like a long one, but shorter!

```
-00a97c0d3868cd856ac6076910ac2c602800;rs=1/2;ss=3.56;t=1639228834.048
```

You can use the `--apd` (ADS-B partial decode) flag to add a partial decode to the comments at the end of an ADS-B message. See the documentation for `ec_978.py` for further information.

Theory of operation

`demod_978`

`demod_978` receives raw FSK data from an SDR radio at the Nyquist limit of twice the bit rate. With a bit rate of 1.041667 Mhz, the sample rate is 2.083334 Mhz. Each sample is a complex IQ value with the real and complex parts being 16 bit integers.

Demodulation is accomplished using the formula:

```
sample = (i[n-2] * q[n]) - (i[n] * q[n-2])  
         -----  
         i^2[n] + q^2[n]
```

where n is the current sample and n-2 is the sample 2 samples before the current sample.

This formula is the equivalent of taking the arctangent and differentiating it for time. It's simple and fast and doesn't require any arctangent tables or arctangent calculations. This technique is from Richard Lyons in *Understanding Digital Signal Processing, Second Edition*. You can find an explanation of this technique [here](#).

The denominator of this equation is for scaling. For our calculations we ignore it. Empirically, you will get slightly more decodes with scaling, but none that can't be corrected in `ec_978.py`.

After demodulating the signal we need to match the sync codes. The sync codes are 36 bit codes and we need to match 32 (or more) out of the 36 bits (32 isn't a magic number— it just represents a reasonable value between too many and too few sync code matches). The sync code for FIS-B is `0x153225b1d` and `0xeacdda4e2` for ADS-B. They are actually inverses of each other, so you could calculate the sync for one, and you would know if the other matched too. Unfortunately, this technique is much slower than using Brian Kernighan's algorithm for calculating 1 bits separately for each sync code. The sync candidate is XOR'ed with the sync word and the one bits counted. If you get more than 4 ones, you can stop— it didn't match.

One quick note: searching for the sync word is very slow using numpy, and is the reason we have a separate program in C. Numpy is quick for all other operations.

Before we even try to match a sync word, we take the additional step of maintaining a 72 bit running total of the absolute values of the samples. Sample values when signal is present are much higher than when only noise is present. In order to even attempt to match a sync word, we must have a value greater than some number. In our case, the default (empirically derived) is 900000. To keep things simpler, all values are presented to the user in millionths. So 900000 is denoted as 0.9. This value doesn't not apply to other SDR setups or amplifications. The `demod_978` program will let the user set this with the `-l` argument. It is probably best to set this to `-l 0.0` and look at the results to find the best level. The `--ll` argument in `ec_978.py` can help with this. Also note that 'fishb_978', with the default level of 0.9, is tuned for FIS-B, not ADS-B. Levels of 0.1 (maybe lower) are required for full ADS-B decoding. The trade off for lower levels is up to a magnitude more garbage noise packets.

Once we have matched a sync code, we will send 8835 32-bit signed integers for a FIS-B packet and 771 32-bit signed integers for an ADS-B packet. These numbers include all the bits required for the message, plus the bits in between the sample bits, plus one extra sample at the beginning and two extra samples at the end. This will allow `ec_978.py` to try some weighted averages to find better sampling points.

The packets are preceded by a 36 character string which tells `ec_978` information about the packet to follow. This includes the type (FIS-B or ADS-B), signal strength, time the packet arrived, and number of mismatched sync bits. The `demod_978.c` documentation contains details on the format of this string. The string is important so that `ec_978.py` will know how many bytes to read for the packet.

We send a single length packet for both ADS-B short and ADS-B long packets. Technically, we could guess at the type since the first five bits of an ADS-B short packet are zero, but we haven't done error correction yet, so we might be wrong.

What we don't do, and might be a future enhancement, is that once we match a sync code, we send the data, and then start looking for the next sync after the end of the packet, not with the next bits. For FIS-B, this isn't an issue, but might be for ADS-B. One case that is quite common is that one set of bits may match a sync code, and the one right next to it (i.e. the 'other' sample in 'every other sample') will match too. We take care of this by sending enough bits in the packet so `ec_978.py` can check the current sample, as well as the sample right after it.

The last thing `demod_978.py` does is to send the 36 character string and packet.

A couple of caveats. This program is written for speed. It uses type-punning to convert between bytes and various size integers. It needs a compiler that allows this, such as GCC. It also assumes little-endian architectures.

ec_978.py

`ec_978.py` receives the fixed length string and reads the appropriate number of bytes for the actual packet. It then turns this into a numpy array. This array is processed slightly differently for ADS-B and FIS-B because FIS-B packets contain six different error correction blocks. I will explain the process for FIS-B packets because ADS-B packets are just a subset.

We take each block of a FIS-B message and try to apply Reed-Solomon error correction to it. We do this by taking the packet and turning it into three packets. One packet is the original packet, one is the set of bits before each bit of the current packet, and one is the set of bits after the current packet. The routine that does this also handles deinterleaving the blocks. This applies only to FIS-B. FIS-B packets are interleaved to help minimize the effect of burst errors.

The first task is to try to decode the original packet without any help from the bits before or after. This works most of the time. But if that doesn't work, we switch to method two.

If you are sampling at many times the bit rate, there is a good chance that one of your samples is close to optimum. When you are sampling at the Nyquist limit of 2 samples per bit at twice the bit frequency, there is a good chance that neither of your samples are near optimum. Both are probably some shade of 'meh'.

If you have a very strong signal, that means that the one and zero points are widely separated and almost any sampling point will work.

If you don't have a strong signal, the one and zero points are closer together and you will run into problems if the sampling points are far off from optimum.

What we do is to use the bits-before and the bits-after to create essentially a weighted average. We do this for the entire packet and then try to error correct again. For example, assume we are using the bits-after at a level of 90%. This means we take each bit in bit-after, multiply it by 0.9, and add it to the corresponding bit in the original sample then divide by two. After we do this for all bits, we try to error correct again. At any given time, we are using either bits-before or bits-after and a fixed percentage to calculate a new packet. We are essentially nudging the sample bits toward either bits-before or bits-after to find a better sampling point.

After lots of experiments, a table was derived ordering the percentages and whether they are bits-before or bits-after in an order which will decode a packet the quickest. For FIS-B, if we decode a packet at a particular shift level, we will start with that shift level for the next block.

There are other techniques we could do, but are not currently using (mostly because what we do now works and is fast enough) such as using zero crossing to estimate a guess on the percentage to use.

It is very uncommon that it takes more than two or three attempts to decode a packet if the packet is going to decode at all.

If we didn't decode the block, we repeat the process by using the next set of bits. In other words, bits-after becomes the current bits, the current bits become the bits-before, and the bits after the original bits-after becomes the new bits-after. This will result in a small number of additional decodes.

Anytime we decode block 0, we check it to see if it is an empty packet, or it ends somewhere in block 0. If it does, we are done and can just fill all the other blocks with zero.

If we fail decoding, we call this same routine to check for early packet ending, but for blocks beyond block 0. This checking doesn't apply to ADS-B.

If we decode all six blocks we create an output string and send it to standard output.

If the packet doesn't decode, we then try some other techniques. Before talking about them, lets discuss what we don't do.

Much effort was put into a number of techniques that didn't work out. The first of these is *erasures*. Erasures are the siren song of Reed-Solomon. They allow you to specify what parts of the message you think are bad. In FIS-B, Reed-Solomon will find and correct up to 10 errors, or detect up to 20 errors (without correction). This is called the Singleton Bound. You can signal a Reed-Solomon decoder to ignore up to 20 bytes that you think are in error. These bytes are the erasures. Typically, if you specify something like 6 erasures, that leaves $(20 - 6) / 2$, or 7 errors that Reed-Solomon can still detect and correct. So the thought is to find the values closest to zero and declare those as erasures. If your message has less than 20 errors, error correction and erasure specification will work great. But if your message has more than this, all bets are off. Adding erasures to a message with more errors than the Singleton Bound will often return a corrected message that is total garbage. In fact, if you add enough erasures, you can get most anything to decode. But it won't decode correctly.

In our case, we would only try to resort to erasures when the initial decode failed at >10 errors. We are then betting that we have 20 or less errors and can replace some of the '*find and correct automatically bytes*' with '*erasure bytes*' (there is 1:2 ratio). And if the message did have less than 20 total errors, this would work. But we have no idea how many errors a message has. And we don't have a good way to check a repaired packet for accuracy. So we abandoned using erasures.

Another technique we tried was expanding on our technique of shifting the bits back and forth to find a better sampling point. You can think of this as horizontal shifting. We added vertical shifting and combined the two. Vertical shifts can nudge bits around zero up or down. Looking at graphs of signals you see a number of cases where this is a problem. In the end, we had the same problem: we could get corrections but couldn't verify their correctness. If FIS-B had an alternative check, like CRC, we could probably get a much higher success rate because we could double check with another method. But it doesn't, so we can't.

So what to do?

We can't use erasures, but we can change data to what we know is true (or almost certain to be true). The easiest example of this is block 0 of FIS-B. The first 8 bytes has a number of bits that are always the same. So we can change the packet to reflect this before we try to error correct it. Similarly, if we are a fixed station and only getting one, or at most a few ground stations, we know what the first six bytes will be. So we can try those. The '`--f6b`' flag in `ec_978.py` will let us set these values. Often, these two techniques alone will allow us to decode FIS-B block 0. Many FIS-B messages only have actual content in block 0 (i.e. they are short messages and the rest of the message are zeros). So by decoding block 0 of a short message, we have avoided having to decode the other blocks.

Another technique is to detect running zeros at the end of a block. Many messages end early, and the rest of the message is only zeros. Block 0 is the best example of this, but it can happen in any block. It's not difficult to come up with a heuristic to find a block with running zeros at the end. If we find running zeros, we set the them to the average zero for the entire block and attempt to decode again. In a noisy packet, some of the zero bits will actually register as one. By forcing these to zero, we may be able to decode the packet.

Together, the above techniques can correct more than 13% of packet decode errors.

It is possible to think of other methods, such as correcting errors in text messages, or using recent past messages to suggest current values (FIS-B likes to send the same things over and over). You don't have to fix many bytes to increase the decode rate. You just need to get under the Singleton Bound.

server_978.py

Nothing fancy here. Just takes standard input and sends it to any connected socket. It is send only. The only wrinkle is that we use `select()` not only for sockets, but also for standard input. This might not work on native Windows, but most likely would work with *Windows Subsystem for Linux*.

Individual program usage

demod_978

`demod_978` reads raw SDR I/Q data from standard input at a frequency of 978Mhz. It assumes 2 samples per data bit, or 2,083,334 samples/sec. Samples should be complex int 16 (sometimes denoted as CS16, or CS16H).

SDR samples are demodulated into packets of signed 32-bit integers. Attributes of each packet (whether FIS-B or ADS-B, arrival time, and signal strength) are stored in a string and sent to standard output. This is followed up with the actual packet data as signed 32-bit integers. These values are then received and processed by the standard input of `ec_978.py`.

The decoding is divided between two programs since searching for sync words in a large amount of data isn't what numpy is best at, but C is amazingly fast at this. Likewise, python, using numpy, is super fast at decoding data packets and trying various approaches to decode data that is at the Nyquist limit.

```
usage: <sdr-program 2083334 CS16> | demod_978 <arguments>
```

Read samples **from SDR and** capture FIS-B **and** ADS-B packets.

Arguments:

`-a`

Process ADS-B packets only. If neither `-a` **or** `-f` are specified, both ADS-B **and** FIS-B are processed. You cannot specify both `-a` **and** `-f` at the same time.

`-f`

Process FIS-B packets only. If neither `-a` **or** `-f` are specified, both ADS-B **and** FIS-B are processed. You cannot specify both `-a` **and** `-f` at the same time.

`-l <float>`

Set the noise cutoff level. Data samples are stronger than the baseline noise level. This sets the minimum value required that demod will attempt to process a packet. The default **is 0.9**. The purpose of this **is** to decrease the number of false packets that are extracted **from noise**. If you are **not** sure **if** you are capturing **all** valid packets, **set** this to **0.0**. The default value has no units, it was determined by evaluation of empirical data. It may vary

based on SDR radio used, **or** SDR program used. Optional.

-x

If you are testing by feeding a file of already captured raw data **in** a file, **set** this argument. 'demod_978' attempts to get the correct timing when a packet arrived, so will figure out how many microseconds past the time the sample was read to provide a correct value. This works fine **for** real-time data, but when dumping a file, it won't work. The -x argument will make sure the times on the packet filename will sort correctly **and** make sense. Optional.

ec_978.py

```
usage: ec_978.py [-h] [--ff] [--fa] [--ll] [--nobzfb] [--noftz] [--apd]
                  [--f6b F6B] [--se SE] [--re RE]
```

ec_978.py: Error correct FIS-B **and** ADS-B demodulated data **from** 'demod_978'.

Accepts FIS-B **and** ADS-B data supplied by 'demod_978' **and** send any output to standard output.

ff, fa

=====

By default will **not** produce any error messages **for** bad packets. To show errored packets use '--ff' **for** FIS-B **and** '--fa' **for** ADS-B.

se, re

=====

The '--se' argument requires a directory where errors will be stored. This directory should be used by the '--re' argument **in** a future run to reprocess the errors. When the '--se' argument **is** given, you need to supply either '--fa', '--ff' **or** both to indicate the **type** of error(s) you wish to save.

When errors are reprocessed **with** '--re', the '--ff' **and** '--fa' arguments are automatically **set**, **and** any '--se' argument **is** ignored.

nobzfb, noftz

=====

Normal operation **is** to attempt to correct known bad FIS-B packets. There are two processes that are used **if** the packet **is not** decoded correctly. The first **is** to apply fixed bits (bits that are always **1 or 0**) to the first portion of a message. The other **is** to attempt to find runs of trailing zeros. You can turn these behaviors off.

'--nobzfb' will turn off 'block zero fixed bits'. These are bits **in** FIS-B block zero that have known fixed values (always **1 or** always **0**).

'--noftz' will prevent the recognition of a block **with** trailing zeros (a string of zeros at the end).

f6b

==

If you have a fixed station **and** only receive one, **or** a few ground stations, you can use the '--f6b' (first six **bytes**) flag to force those values **in** packets that initially fail decoding. More than one value can

be listed **if** enclosed **in** quotes **and** separated by spaces. Examples of use would be:

```
--f6b 3514c952d65c  
--f6b '3514c952d65c 38f18185534c'
```

ll

==

demod_978 uses a cutoff signal level to avoid trying to decode noise packets that have the correct sync bits. '**--ll**' (lowest level) will display on standard error the lowest level that either a FIS-B packet **or** ADS-B packet (each have their own lowest level) decoded at. If a lower level **is** found, it **is** displayed. This **is** a good way to find an optimal setting **for** demod_978's '**-l**' switch.

apd

==

This stands **for** ADS-B partial decode **and** will add an additional comment to the decode of ADS-B packets.

Example comments are:: :

```
1.2.A79B5F/1.N59DF/8500/G  
0.0.A38101//2275/A  
1.0.A38101/1.1200/2325/A23:32:L7  
2.0.A38101//2275/A05:29-30:M11
```

The first number **is** the payload **type** code. The second **is** the address qualifier (see D0-282B **for** what these values mean).

Then the ICAO aircraft **id** (**or** something that stands **in** **for** it).

The first item within the first pair of slashes **is** the emitter category (lots of options here, but most often: 0=unknown, 1=light acft, 2=small acft, 3-6=heavy acft, 7=heli) **and** the callsign (usually the squawk, N-number, **or** flight callsign). This section **is** optional. The altitude **is** **in** the next set of slashes. The last portion **with** be either a 'G' **is** this **is** a TIS-B/ADS-R message sent by a ground station, **or** 'A' **if** a UAT message sent directly **from** **an** aircraft.

If the message **is** of **type** 'A', **and** the aircraft has received **any** FIS-B messages **from** **a** ground station, the aircraft can show how many messages it has received **from** **a** particular ground station on a particular data channel. Each ground station has a 'TIS-B Site ID' **in** the **range** of 1 to 15. Each site **id** **is** allocated a particular **set** of channel numbers that it will transmit on. High power stations get 4 channels, medium stations 3, low power stations 2, **and** surface stations 1. Each second, there **is** (**for** lack of a better term) the 'data channel of the second'. **This is determined by the number of seconds after UTC midnight.** All ADS-B messages sent by aircraft will send the number of FIS-B packets they received **from** **the** 'data channel of the second' **in** the last 32 seconds. In the comment, this can look like:

'A05:29-30:M11'. 'A' means sent by aircraft, '05' **is** the 'data channel of the second', '29-30' **is** the **range** of FIS-B packets received by the aircraft on that data channel **in** the last 32 seconds, **and** 'M11' maps the data channel back to the TIS-B site **id** **and** the power of the station ('H', 'M', 'L', 'S').

The full data channel FIS-B packets received section **is not** sent **if** the number of packets received **is** zero. Some planes never report **any** packets.

```
fet
===
Display extra timing information in the FIS-B 'rs=' string. Using this
from a playback file will not give accurate results. Use this only with
real-time data. A typical string would look like:
```

```
/28:594:154.5/75/12/4
```

The first three numbers (28:594:154.5) are separated by colons because they are actually three different manifestation of the same thing. The first (28) **is** the 'Transmission Time Slot' (1-32) which relates to the actual time the message was sent. '594' **is** the 'Message Start Opportunity' or 'MS0'. All data sent in UAT messages (ground messages or ADS-B) are referenced to MSOs. '154.5' **is** the actual millisecond time of the current second the message was sent. There are formulas that can directly convert between the first three values. They **all** represent the same point **in** time.

The **next** number (75) **is** the time difference **in** milliseconds **from when** the message was transmitted to when it was received. This represents the distance **from the** station to the receiver, the time spent **in** the radio, **and** the computer processing time (mostly computer processing time).

'12' **is** the 'TIS-B Site ID' (1-15). This **is** a fixed value **for** each ground station. It defines which data channels will be used.

'4' **is** the 'Data Channel' **for** this transmission (1-32). It **is** directly tied to the TIS-B Site ID.

```
d978, d978fa
=====
```

If you are piping the output of FISB-978 to a program that expects output **from the** original dump-978 **or** Flight Aware's dump-978 programs, you can supply '--d978' **or** '--d978fa' to produce the exact output those programs would produce. Note that this includes less information than FISB-978 normally produces. All ADS-B **and** FIS-B errors are also suppressed. It **is** also an error to use both switches at once.

```
saveraw
=====
```

Saves the raw demod_978 output **in** a disk file. This **is** usually done briefly **for** later analysis (such **as for** eye diagrams). Each packet **is** save **in** its own file. The files are saved **in** the current directory **and** have a name of the form: '1646349680.227.F.i32' where '1646349680.227' **is** the UTC epoch time of arrival, 'F' means FIS-B ('A' means ADS-B, either short **or** long). The extension **is** always '.i32'.

Optional Arguments

```
-----
```

- h, --help show this help message **and** exit
- ff Print failed FIS-B packet information **as** a comment.
- fa Print failed ADS-B packet information **as** a comment.
- ll Print lowest levels of FIS-B **and** ADS-B signal levels.
- nobzfb Don't repair block zero fixed bits.
- noftz Don't fix trailing zeros.
- apd Do a partial decode of ADS-B messages.
- fet Show FIS-B extra timing information.
- f6b F6B Hex strings of first **6 bytes** of block zero.
- se SE Directory to save failed error corrections.

```
--re RE      Directory to reprocess errors.  
--d978       Mimic dump978 output format.  
--d978fa     Mimic dump978-fa output format.  
--saveraw    Save demod_978 output in file.
```

server_978.py

`server_978.py` just takes the standard output from `ec_978.py` and serves it to whoever connects. It is run as

```
./server.py --port 3333
```

The default port is 3333 so you can omit the port argument if that is the one you want.

eye.py

`eye.py` is a program that will take demodulated data from `ec_978.py` generated by using the `--saveraw` command and create eye diagrams in a number of forms.

See the instructions above and make sure that `numpy`, `matplotlib`, `scipy`, and `colour` are installed via `pip3`.

The program requires the filenames produced by `ec_978.py` to work correctly. To display all of the available eye diagrams, type something like:

```
./eye.py eye-samples/1647067220.215.F.i32
```

It will then display 6 different eye diagrams, one at a time. You can get to the next one by closing the window with the current diagram—the next one will pop up.

The directory `eye-samples` has some example demodulated data you can experiment with. Examples of images are shown after the following usage instructions.

```
usage: eye.py [-h] [--all] [--adpns] [--adps] [--adpgfb] [--adpgbf] [--flns]  
              [--fls] fname
```

`eye.py`: Show signal **as** eye diagram.

Uses the data `from 'ec_978.py'` **with** the `'--saveraw'` flag **set** to show an eye diagram.

Not specifying **any** optional arguments implies **all** eye diagrams will be displayed **in** sequence.

Each diagram shows approximately **3** data points (**6 + 1** samples). Data **is** sampled at **2.083334** Mhz, **or** twice the data rate of **1.041667** Mhz. Non-smoothed eyes shows only the actual samples. Smoothed samples are up-sampled **and** interpolated by a factor of eight.

Either FIS-B **or** ADSB data can be displayed. ADS-B **is** tuned **for** long ADS-B packets (the most common). If, when looking at an ADS-B eye diagram, you see a thin horizontal line, this denotes a short packet. The horizontal line represents no data transmission which **is** usually at a very low magnitude compared **with** the signal. First **and** last **10%** diagrams will only show the first **10%**, because the last **10%** **is not**

modulated.

Some eye diagrams only show the first 10% of data, and the last 10% of data. For the non-smoothed version, the first 10% is shown in gray, and the last 10% in tan. This way you can detect an abnormal sampling frequency. The two colors should mostly overlap.

The smoothed first and last 10% uses a normal color scheme. A full FIS-B packet can clutter up the display, and only showing 20% of samples is easier to interpret. Also, FIS-B packets have ones and zeroes at both ends, while placeholder packets have only zeros in the middle.

Data points can be shown with 'gradient colors'. Early data will start off in red, then gradually turn to yellow, green, and finally blue. You can use gradient colors to observe outlier values and tell where within the packet they occurred. '--adpgfb' displays the data front to back (red to blue), and '--adpgbf' will display the data back to front (blue to red). Red will always be the early data and blue the last data. Gradient colors also help visualize the eye better in noisy signals.

If you are displaying more than one diagram, the next diagram will appear after closing the window of the current diagram.

positional arguments:

 fname Filename to use.

optional arguments:

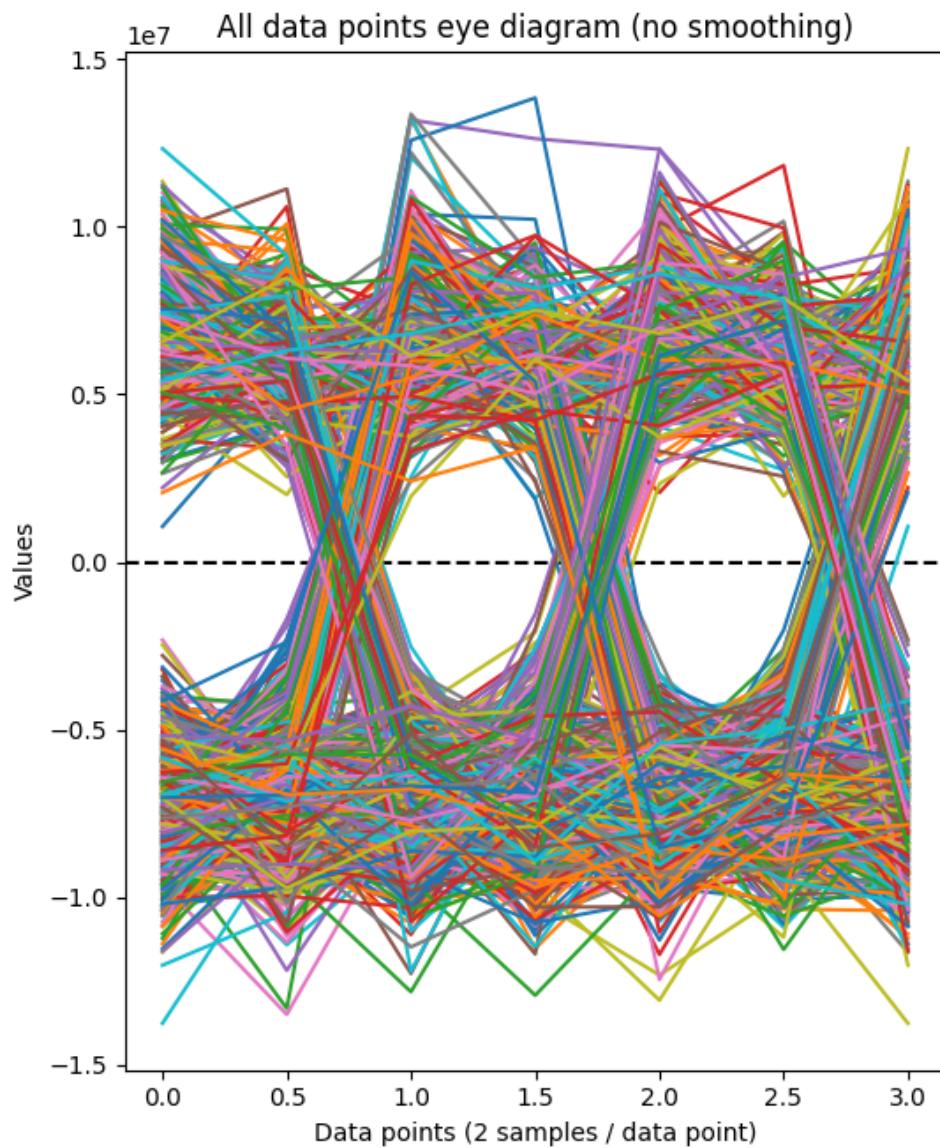
- h, --help show this help message and exit
- all Display all graphs in sequence.
- adpns Show all data points no smoothing.
- adps Show all data points with smoothing.
- adpgfb Show all data points gradient colors front to back.
- adpgbf Show all data points gradient colors back to front.
- flns Show first and last 10 pct no smoothing.
- fls Show first and last 10 pct with smoothing.

Here are examples of the the various eye diagrams available. These represent a very strong signal received from about 4 miles away through a number of obstacles.

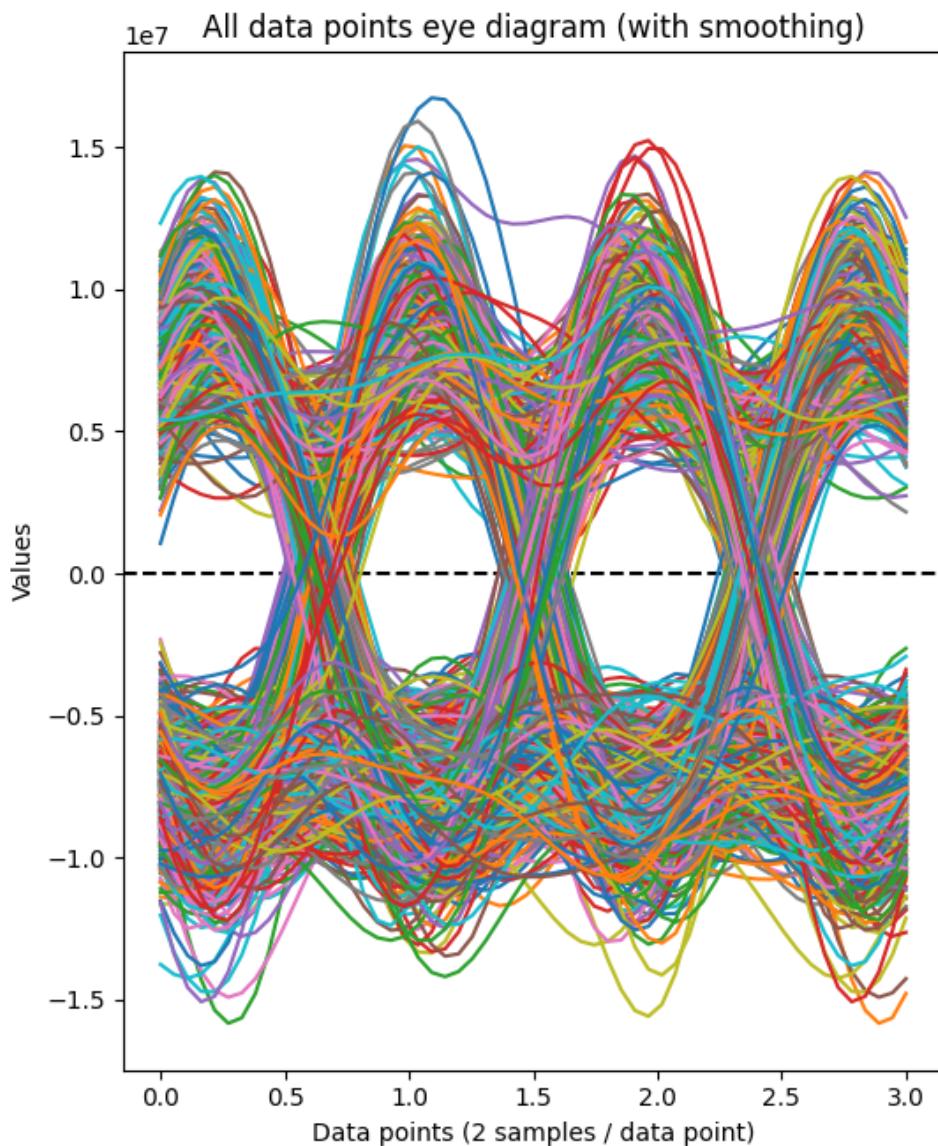
Each eye diagram is composed of a number of overlayed samples. In our case, we collect 7 samples that are displayed on the x-axis with the values 0 to 3. The magnitude of each point is the y-axis. Starting with the last bit we used, we collect the next 6 samples and overlay that over the first sample. We continue to do that until we run out of samples. The result is a complete eye diagram.

The type of eye diagram is denoted by the command flag to display it. You can specify as many eye diagram types as you wish. Using the --all flag will display all available types, as will not including any flags.

--adpns (all data points, no smoothing) Shows all data points exactly as they are demodulated. Since the sampling rate is twice the data bit frequency, each data bit only consists of two samples. Don't forget, that while we draw a line between the sample points (that's what eye diagrams are all about), we actually don't know what the actual data between those points is.



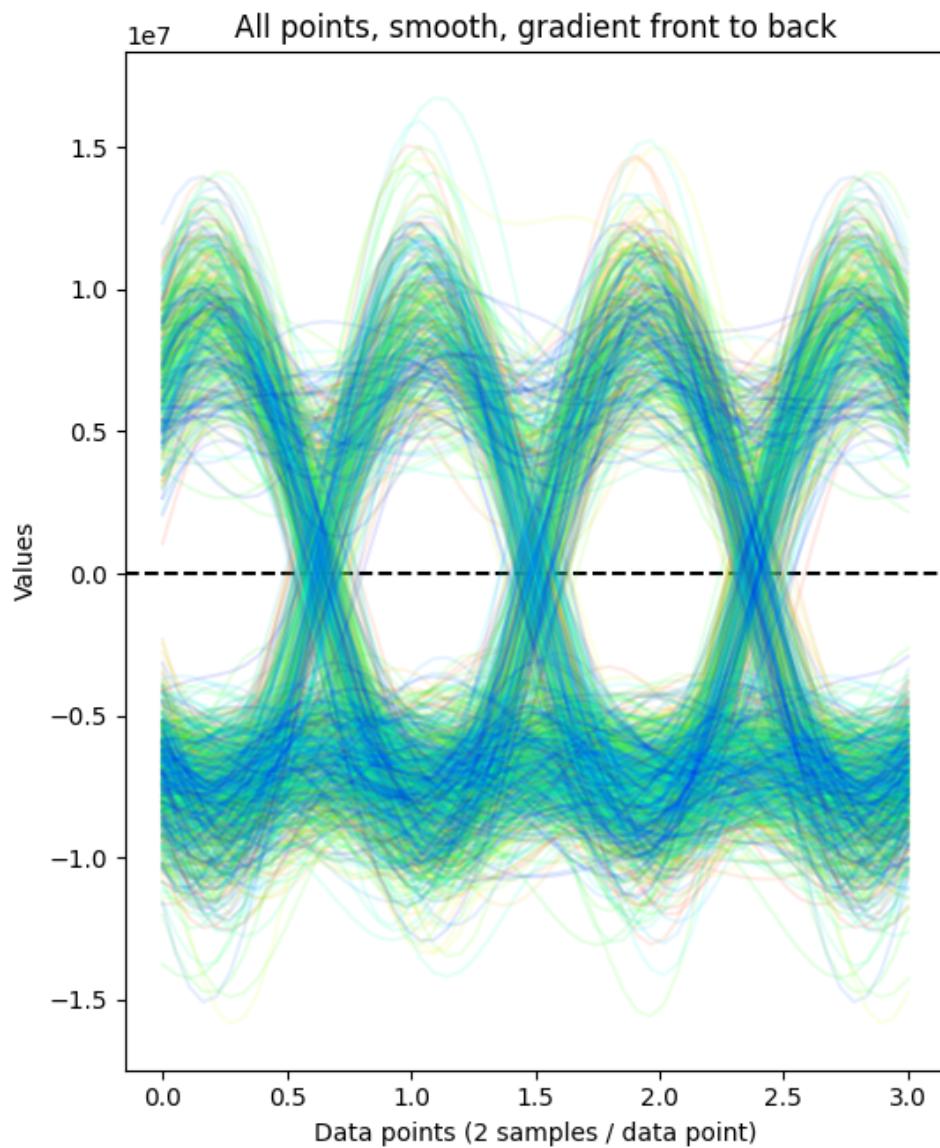
--adps (all data points, smoothing) Same as above, but the data is interpolated by 8 to present a smooth curve. Remember, the points between the actual samples do not represent the truth, only an estimate.



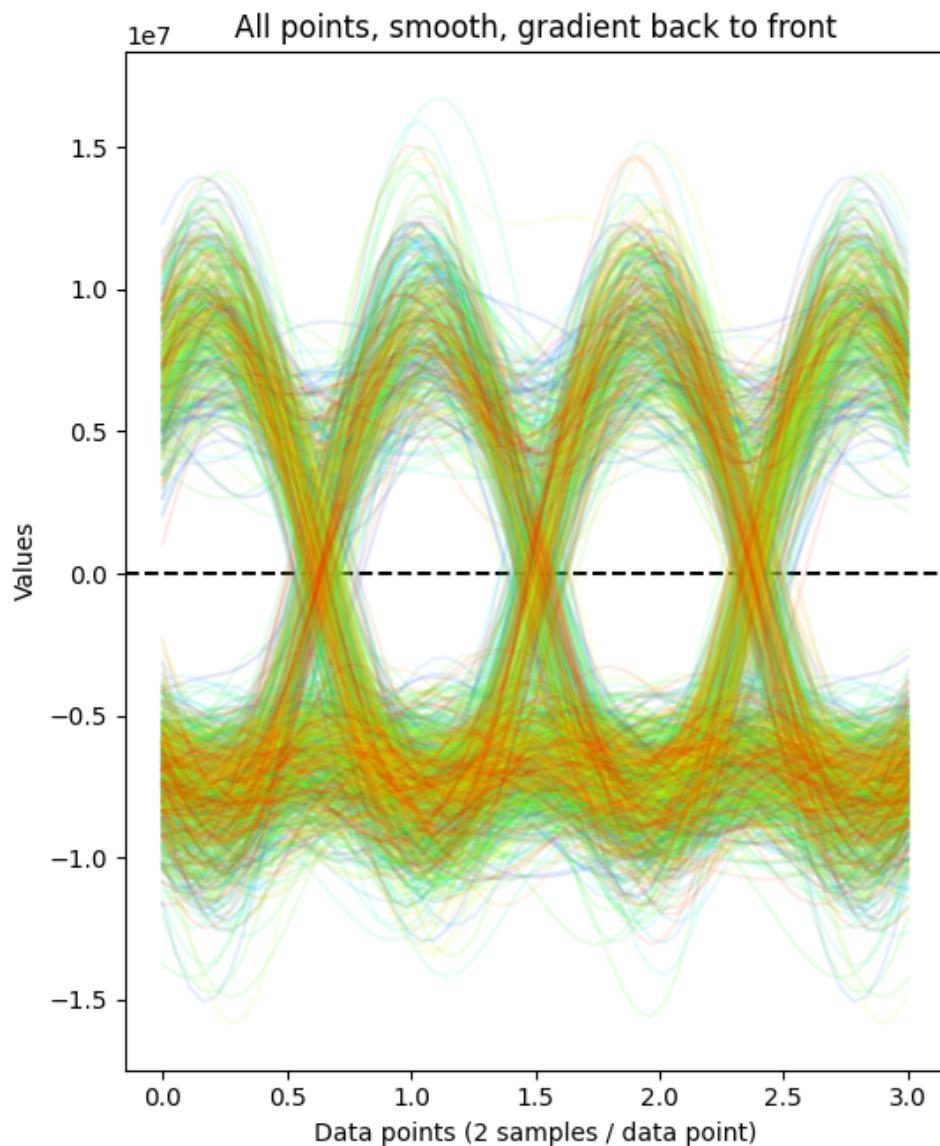
`--adpgfb` (all data points, gradient, front to back) This is like the previous diagram, but instead of using a different color for each plot, we use a continuum of colors. The first part of the packet starts off red, then gradually turns to yellow, then green, and finally blue. Also, the transparency is changed to translucent, so you can see some of the plots that have been overlayed. Because the color blue represents the last of the plots, these eye-diagrams tend to hide the red and yellow colors.

By looking for outlier samples, you can tell what part of the packet they came from. This can help in identifying burst errors.

These plots tend to emphasize data density since places where more data occurs shows up denser. Outliers are still visible, but fainter. If you are looking at a noisy signal, where there doesn't appear to be an eye at all on the default plot, this plot can often delineate the eye.

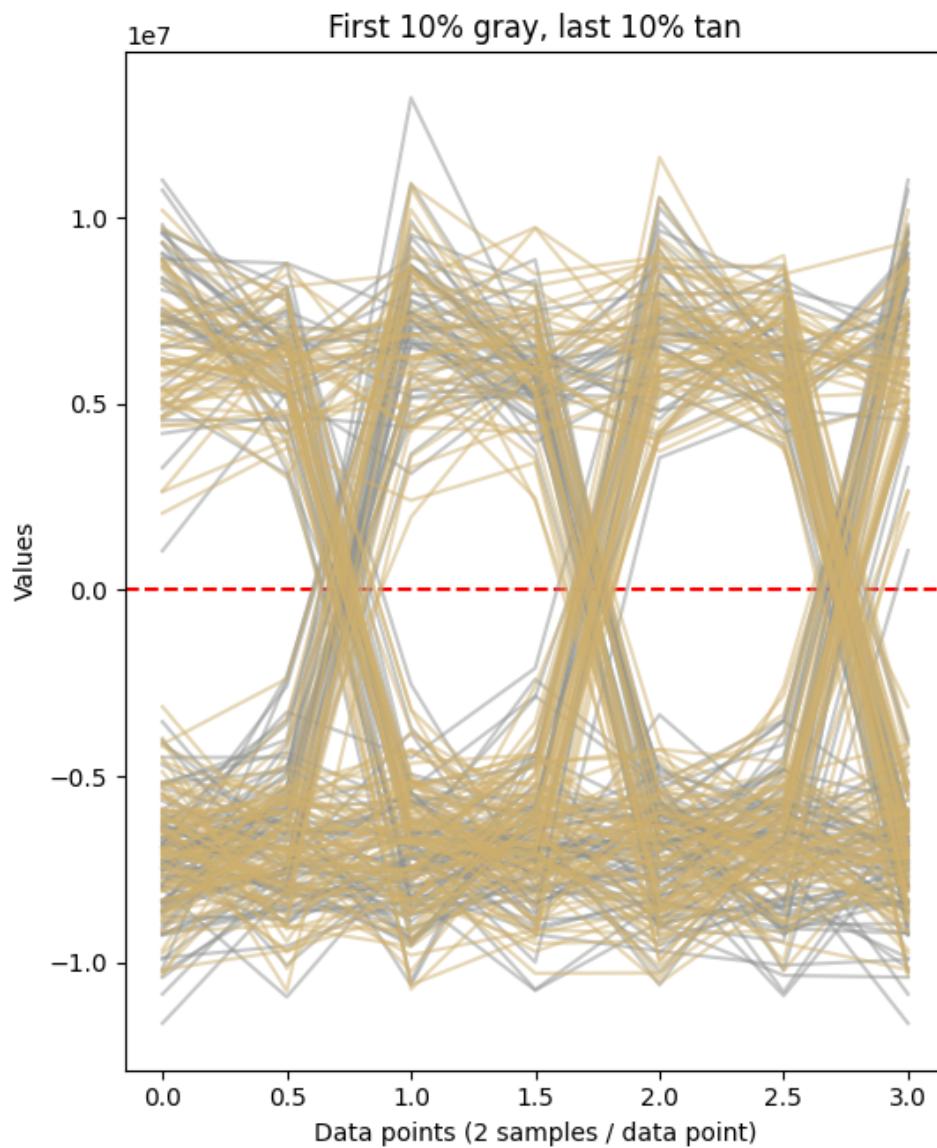


--adpgbf (all data points, gradient, back to front) This is exactly the same as the previous plot, except the last packet data is plotted first and the first data is plotted last. Red still represents early data and blue is late data. These will appear more reddish, since the red data is plotted last.

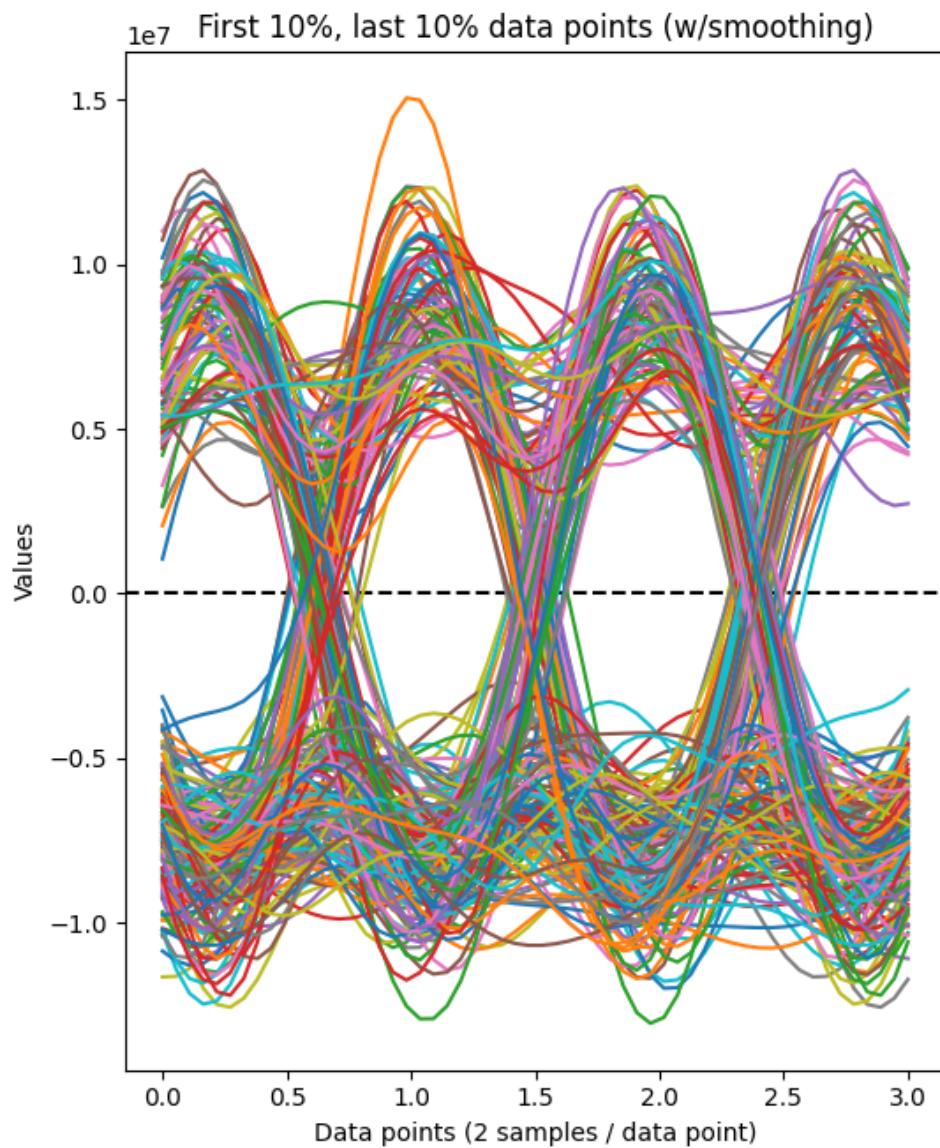


--flns (first/last, no smoothing) First and last plots show only the first 10% of the data and the last 10% of the data. All FIS-b packets have some 1's and 0's at the front and the at the end (parity). In the middle, they can be all zeroes (place-keeping packets).

This particular eye diagram uses gray for the early data and tan for the last data. There is also some transparency applied. The main purpose is to look for shifting from the beginning to the end of the packet, indicating the sampling rate might be off.

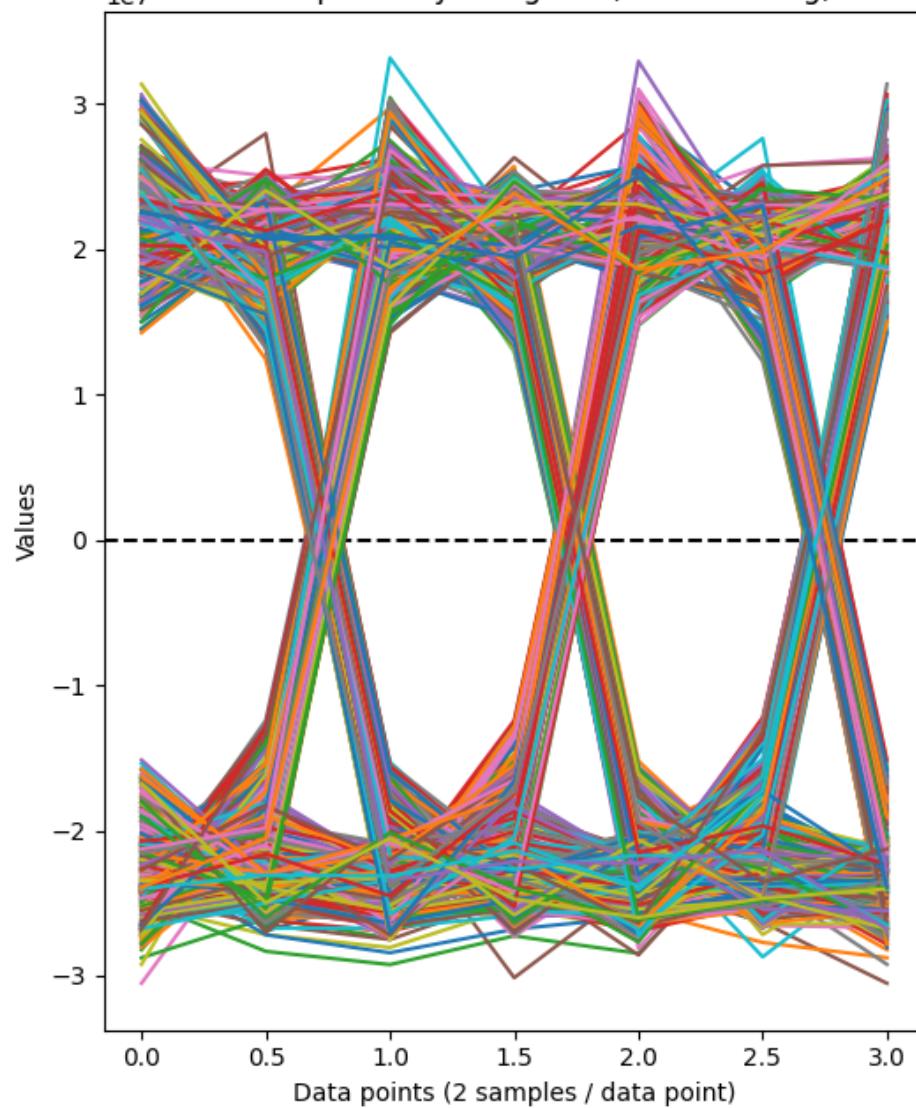


--fls (first/last, smoothing) Same as the last diagram, except smooths the packets and uses a different color for each data plot and no transparency (same as --adpns). If you have a very dense and noisy diagram, this can help by only showing 20% of the data.

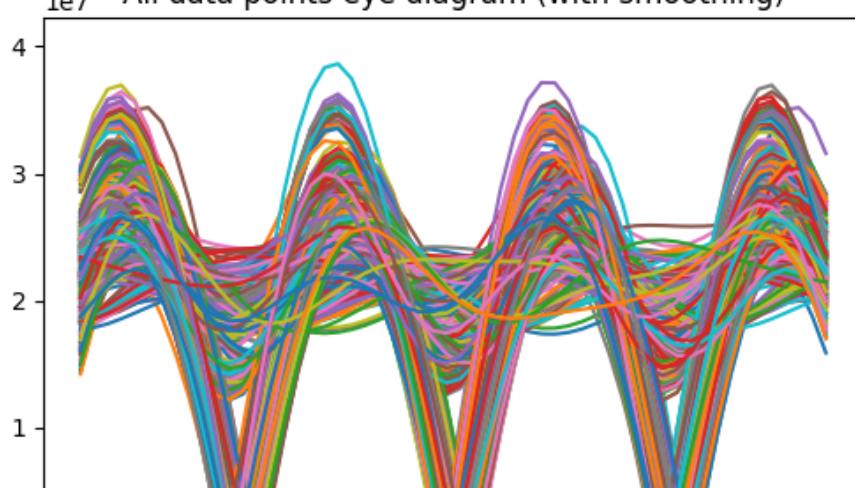


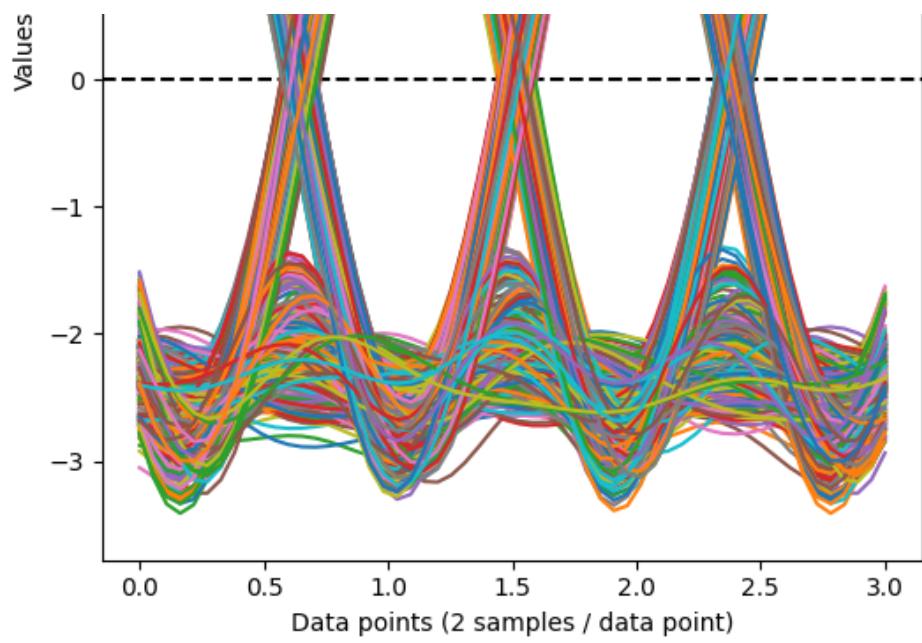
For reference purposes, here are examples from a FIS-B transmission received a few hundred feet from a tower. These represent the 'gold standard' of reception.

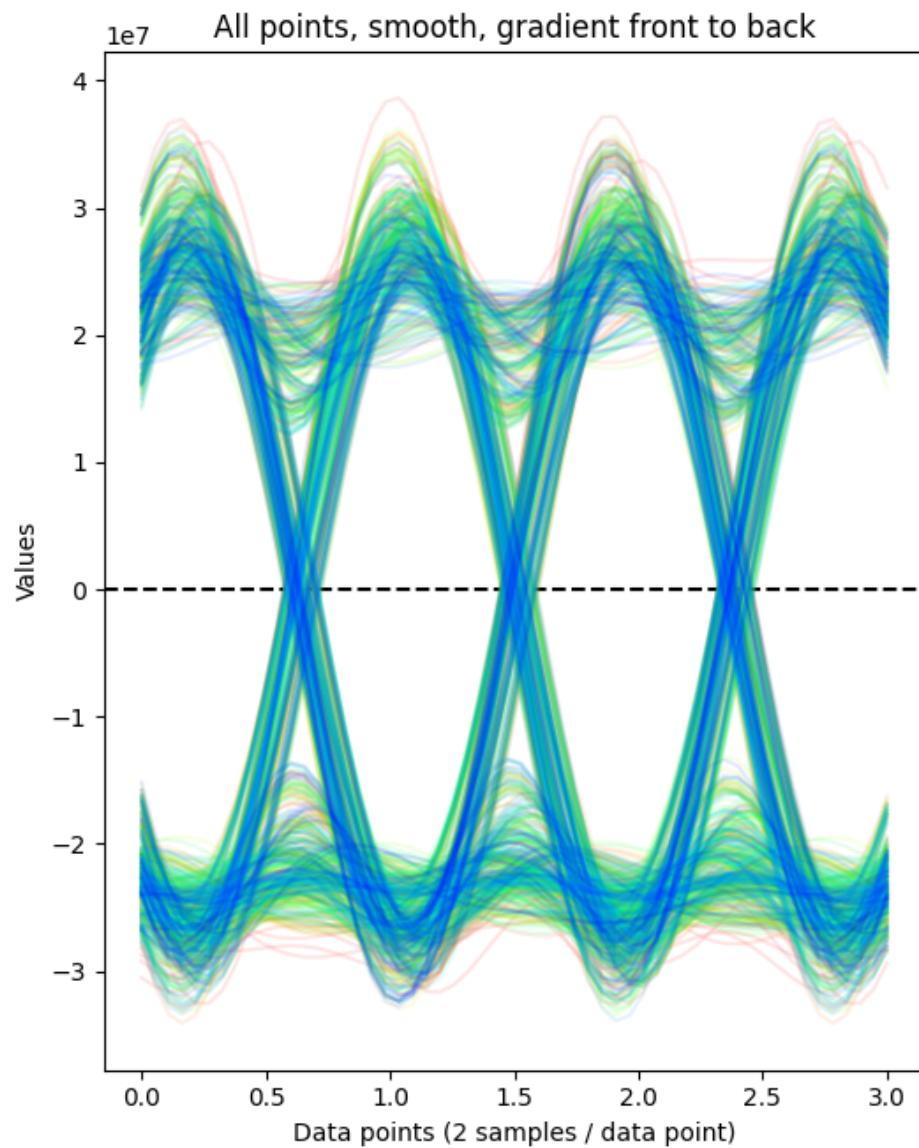
$1e7$ All data points eye diagram (no smoothing)

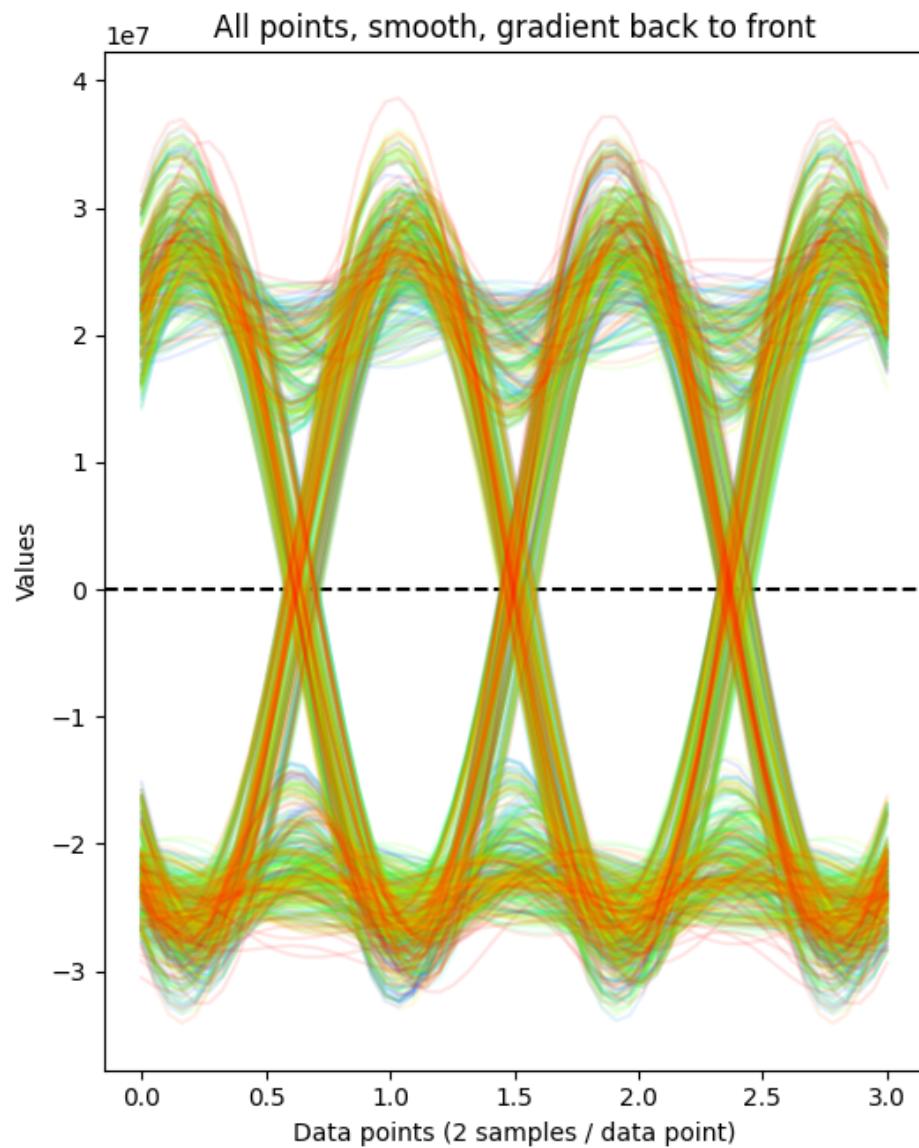


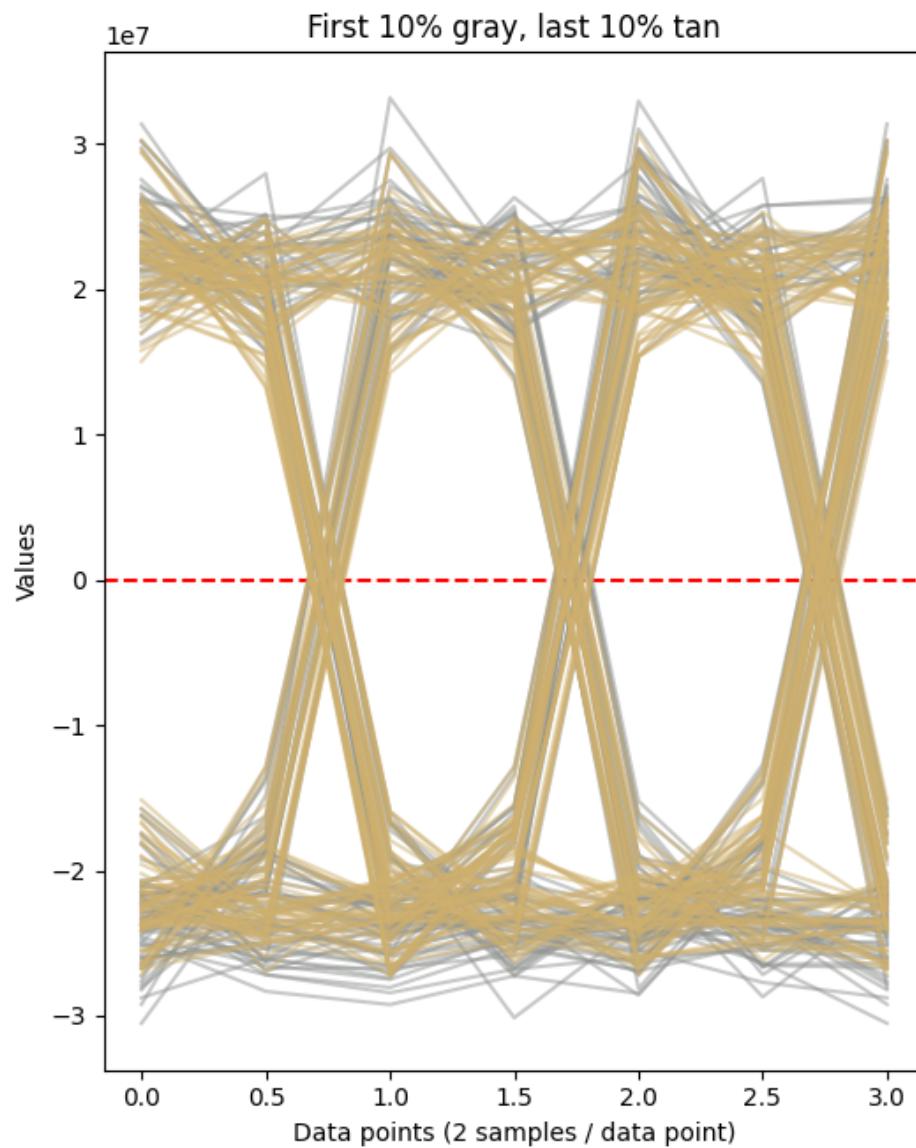
$1e7$ All data points eye diagram (with smoothing)

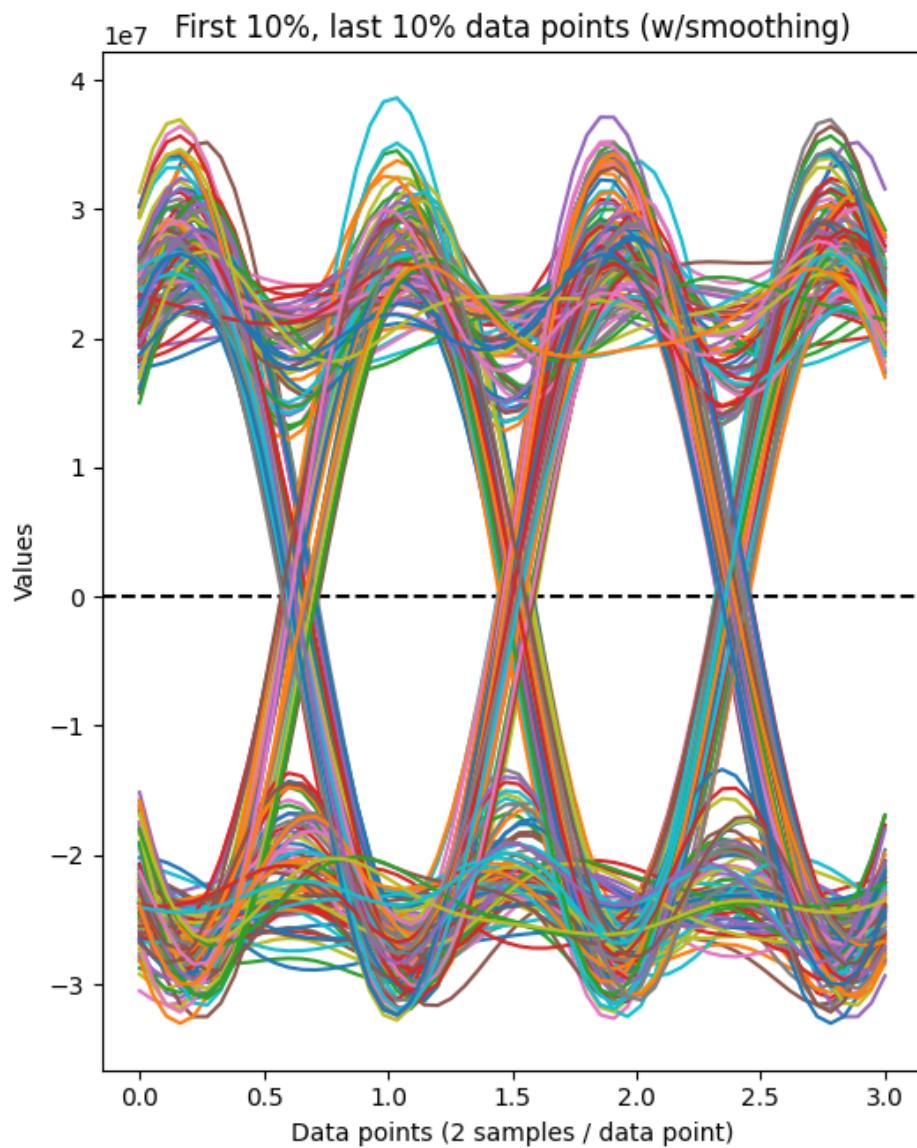












Building Documentation

If you want to build the documentation, install [sphinx](#). On Ubuntu 20.04 you can do this with:

```
sudo apt install python3-sphinx
```

Next, install the Python requirements from the `bin` directory as:

```
pip3 install -r ./misc/requirements-sphinx.txt
```

Next come the tricky parts. You will need to install [doxygen](#) and [breathe](#) on your system. The setup for these is system dependent, so you need to read and follow the documentation.

At a minimum, you will need to edit the file `docs/source/conf.py` and edit the line with `breathe_projects` and change the provided path to reflect where `doxygen/xml` will be on your system. This will normally be the

path to where you cloned fisb-978 (i.e. `~/fisb-978/doxygen/xml`). The `xml` directory won't exist at first, but will be created the first time you build the documentation. There is a doxygen project file in `doxygen/fisb_978`.

Then (assuming 'fisb-978' was cloned in your home directory):

```
cd ~/fisb-978/docs  
./makedocs
```

The html documentation will be found in `fisb-978/docs/build/html`. Load `index.html` in your browser to view. Sphinx is configured to link directly to the source, so this is an easy way to explore the code.